

REPORT

PROJECT-2

TITLE:Credit Card Behaviour Score Prediction Using Classification and Risk-Based Techniques

Name— PONNAGANTI PAVAN KUMAR

ENROLLMENT NUMBER— 24114065

OVERVIEW ABOUT PROJECT:

Bank A aims to improve its credit risk management framework by developing a forward-looking Behaviour Score — a classification model that predicts whether a credit card customer will default in the following month.

We are provided with anonymized historical behavioral data of over 30,000 credit card customers, with a labeled target variable: default.payment.next.month. This variable indicates whether a customer defaulted on their payment in the next billing cycle. The goal is to build a model that can accurately flag potential defaulters in advance, allowing the bank to adjust credit exposure, trigger early warning systems, and prioritize risk-based actions.

FINAL GOAL:

The goal is to go beyond prediction and create a financially interpretable model that helps the bank understand default patterns, take early action, and manage credit exposure.

PROJECT OBJECTIVES:

Build a binary classification model to predict customer credit card default using behavioral and financial data. Address class imbalance using techniques like SMOTE, class weighting, or downsampling. Conduct in-depth analysis of behavioral trends (e.g., payment delays, utilization, consistency) and engineer financially meaningful features. Evaluate and compare multiple models including Logistic Regression, Decision Trees, and Ensemble methods like XGBoost or LightGBM. Optimize classification threshold based on business risk trade-offs, and select metrics like F1-score, AUC, or Precision to guide model tuning and final prediction. Generate risk-informed predictions on an unlabeled validation set, reflecting real-world credit risk management goals.

VARIABLES:

There are 26 variables:

ID: ID of each client

LIMIT_BAL: Amount of given credit in NT dollars (includes individual and family/supplementary credit)

SEX: Gender (1=male, 2=female)

EDUCATION: (1=graduate school, 2=university, 3=high school, 4=others, 5=unknown, 6=unknown)

MARRIAGE: Marital status (1=married, 2=single, 3=others)

AGE: Age in years

PAY_0: Repayment status in September, 2005 (-1=pay duly, 1=payment delay for one month, 2=payment delay for two months, ... 8=payment delay for eight months, 9=payment delay for nine months and above)

PAY_2: Repayment status in August, 2005 (scale same as above)

PAY_3: Repayment status in July, 2005 (scale same as above)

PAY_4: Repayment status in June, 2005 (scale same as above)

PAY_5: Repayment status in May, 2005 (scale same as above)

PAY_6: Repayment status in April, 2005 (scale same as above)

BILL_AMT1: Amount of bill statement in September, 2005 (NT dollar)

BILL_AMT2: Amount of bill statement in August, 2005 (NT dollar)

BILL_AMT3: Amount of bill statement in July, 2005 (NT dollar)

BILL_AMT4: Amount of bill statement in June, 2005 (NT dollar)

BILL_AMT5: Amount of bill statement in May, 2005 (NT dollar)

BILL_AMT6: Amount of bill statement in April, 2005 (NT dollar)

PAY_AMT1: Amount of previous payment in September, 2005 (NT dollar)

PAY_AMT2: Amount of previous payment in August, 2005 (NT dollar)

PAY_AMT3: Amount of previous payment in July, 2005 (NT dollar)

PAY_AMT4: Amount of previous payment in June, 2005 (NT dollar)

PAY_AMT5: Amount of previous payment in May, 2005 (NT dollar)

PAY_AMT6: Amount of previous payment in April, 2005 (NT dollar)

PAY_TO_BILL_ratio :Payment-to-bill ratio over 6 months

NEXT_MONTH_DEFAULT: Default payment (1=yes, 0=no)

The progress of notebook is organized as follow:

- DATA LOADING AND DATA VISUALIZATION
- Data Preprocessing
- data preparation
- data visualization
- Features Engineering and Features selection
- Model Development
- Model Evaluations

DATA LOADING AND DESCRIPTION :

In this credit default prediction project, we work with two key datasets: **train_dataset_final1.csv** and **validation.csv**.

The **train_dataset_final1.csv** file contains historical data of around 25,000 customers, including both the input features (such as age, credit limit, payment history, etc.) and the target variable indicating whether each customer defaulted or not. This dataset is used to train and evaluate various machine learning models by applying preprocessing steps like missing value imputation, normalization, and categorical encoding.

The second dataset, **validation.csv**, which contains approximately 5,000 unlabeled customer records. This dataset has the same features as the training data but lacks the target variable. The trained model predicts whether each customer in this validation set is likely to default in the next month, and the results are saved in a CSV file containing the **Customer_ID** and the predicted default label. This final prediction file can then be used for business decision-making or submission.

Data Displaying:

This line uses the **read_csv()** function from the Pandas library to load a CSV file containing the training dataset. The file is located at the specified path on your

computer. Once loaded, the data is stored IN the DataFrame `df`, which organizes it into rows and columns for easy analysis and preprocessing in Python.

WHAT WE HAVE DONE IN DATA LOADING AND MISSING VALUES ANALYSIS

In this step, we performed an initial inspection of the dataset to understand its structure and quality. The `'df.info()'` function was used to display key details such as the number of entries, column names, data types, and non-null counts, which helps us assess the overall shape and consistency of the data. Following that, `'df.isnull().sum()'` was executed to check how many missing values exist in each column, allowing us to identify features that may require cleaning or imputation. Finally, `'df.isnull().sum().sum()'` was used to calculate the total number of missing values across the entire dataset, giving us a quick sense of the completeness of the data before moving forward with further analysis.

OBSERVATIONS

We have identified a total of 126 missing values in the training dataset, all of which are located in the AGE column. To ensure the dataset is suitable for machine learning modeling, it is essential to address these missing values during the data cleaning process. Handling null values appropriately is a crucial step, as many ML algorithms do not perform well with incomplete data. Therefore, we will proceed to impute or replace these missing age values using a suitable strategy, such as mean, median, or predictive imputation, depending on the distribution and business context.

Exploratory Data Analysis: Visualization of Risk Drivers

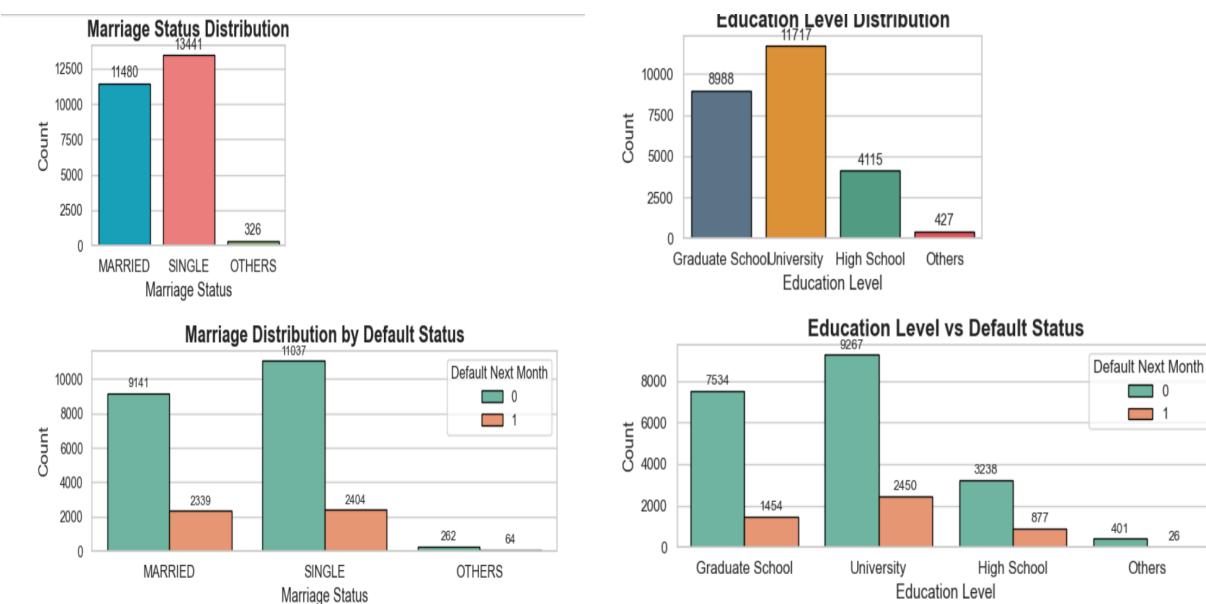
In this section, we perform an in-depth exploratory data analysis (EDA) using rich visualizations to uncover key patterns and relationships relevant to credit card default risk.

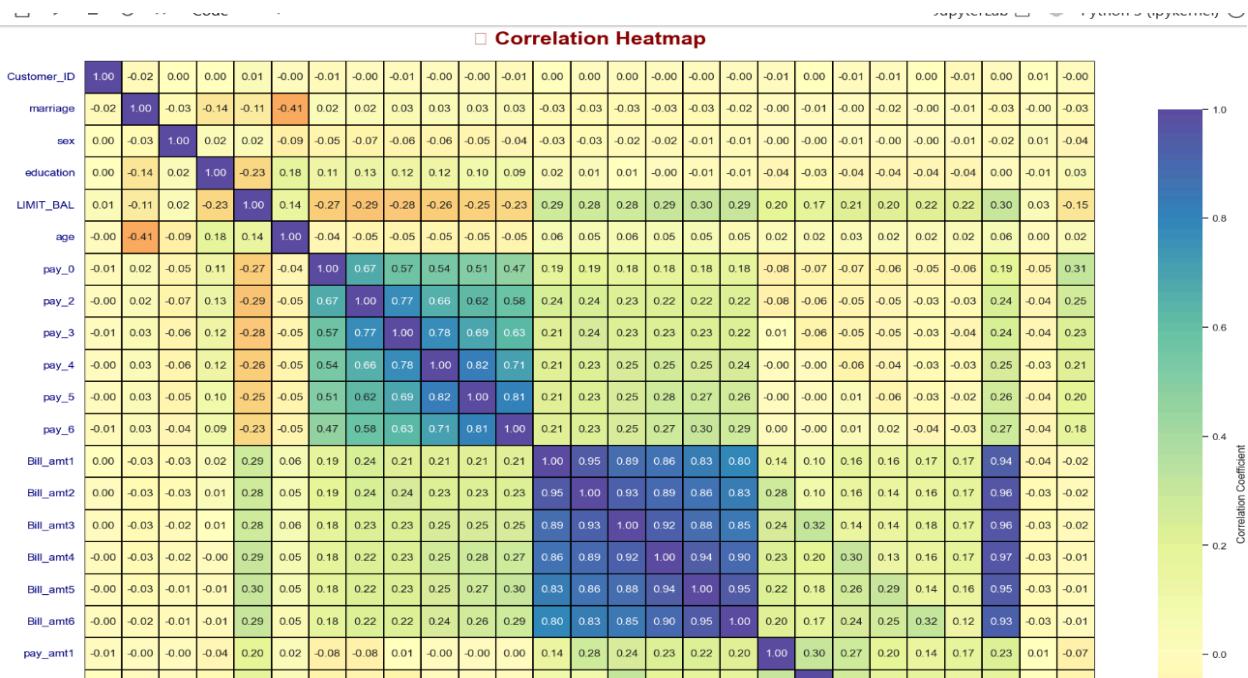
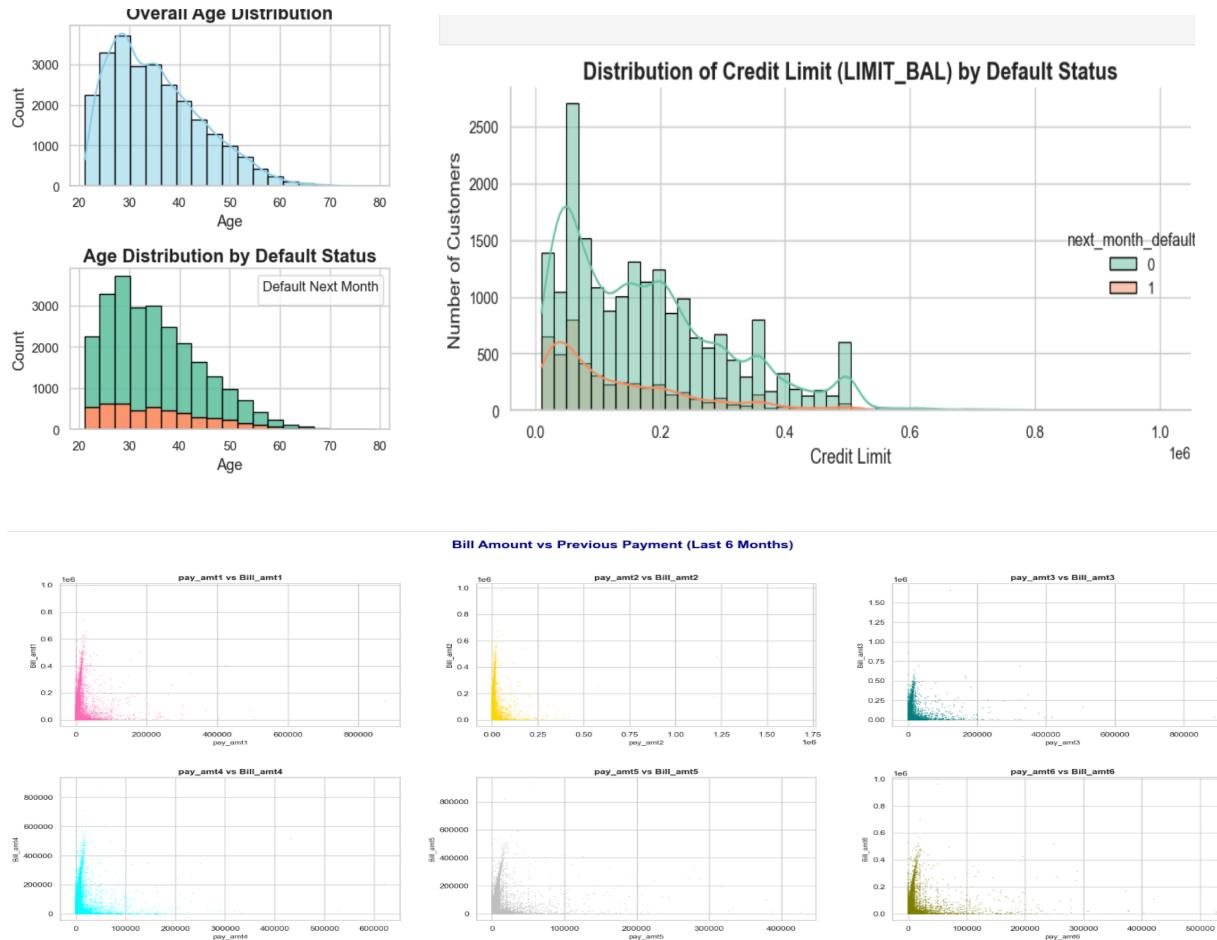
Purpose:

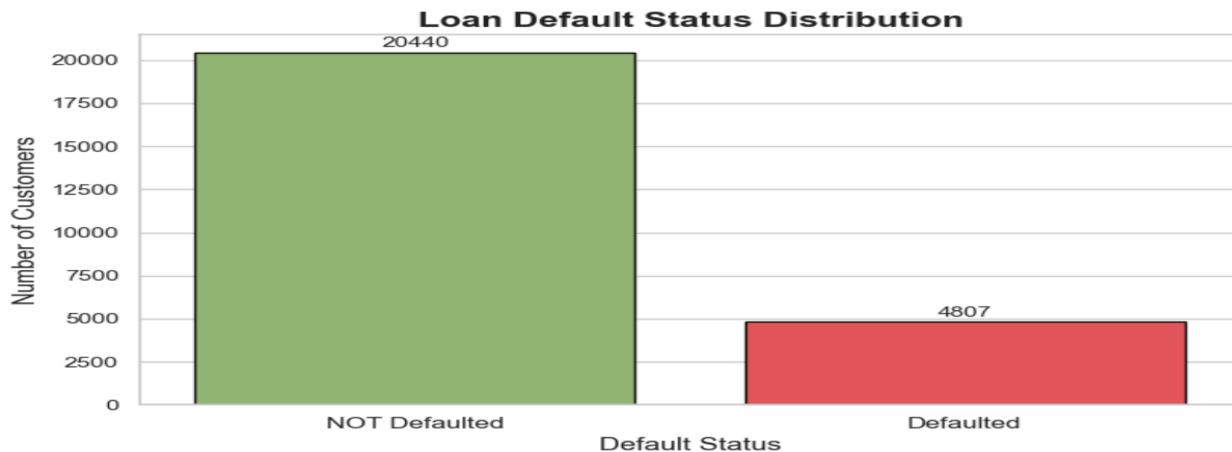
- To gain actionable business and modeling insights by exploring how demographic factors, financial indicators, and payment behaviors differ between defaulters and non-defaulters.
- To visually assess class imbalance, variable distributions, and multicollinearity, informing both feature engineering and model development.

Key Analyses & Visualizations:

- **Target Variable Distribution:** Quantifies the proportion of defaulting vs. non-defaulting customers, highlighting the presence of class imbalance—a critical consideration for credit risk modeling.
- **Demographic Analysis:** Examines how gender, marital status, education, sex, LIMIT_BAL, BILL Amount vs previous payment, age are associated with default risk, helping identify high-risk segments for targeted interventions.
- **Financial Indicators:** Compares distributions of credit limit, average bill amount, payment-to-bill ratio, and bill-to-limit ratio between defaulters and non-defaulters to reveal key financial behaviors linked to credit risk.
- **Payment Patterns:** Visualizes historical payment status across multiple months, both as a heatmap and as trends, to uncover behavioral signatures (such as chronic delinquencies) that precede default events.
- **Correlation Analysis:** Evaluates feature interrelationships and their direct correlation with default, guiding the selection of the most predictive and non-redundant variables for modeling.







The data is quite imbalance which about 19% of clients will default next month.

Advanced Exploratory Data Analysis: Financial & Behavioral Insights

In this section, we build on initial EDA by using advanced statistical visualizations and domain-driven metrics to deepen our understanding of default risk in credit card customers

Purpose:

- To visualize and interpret how engineered financial ratios, payment patterns, and demographic segments relate to default risk.
- To surface actionable business insights and validate which behavioral and financial features are most predictive for modeling.

Average Payment Delay Status

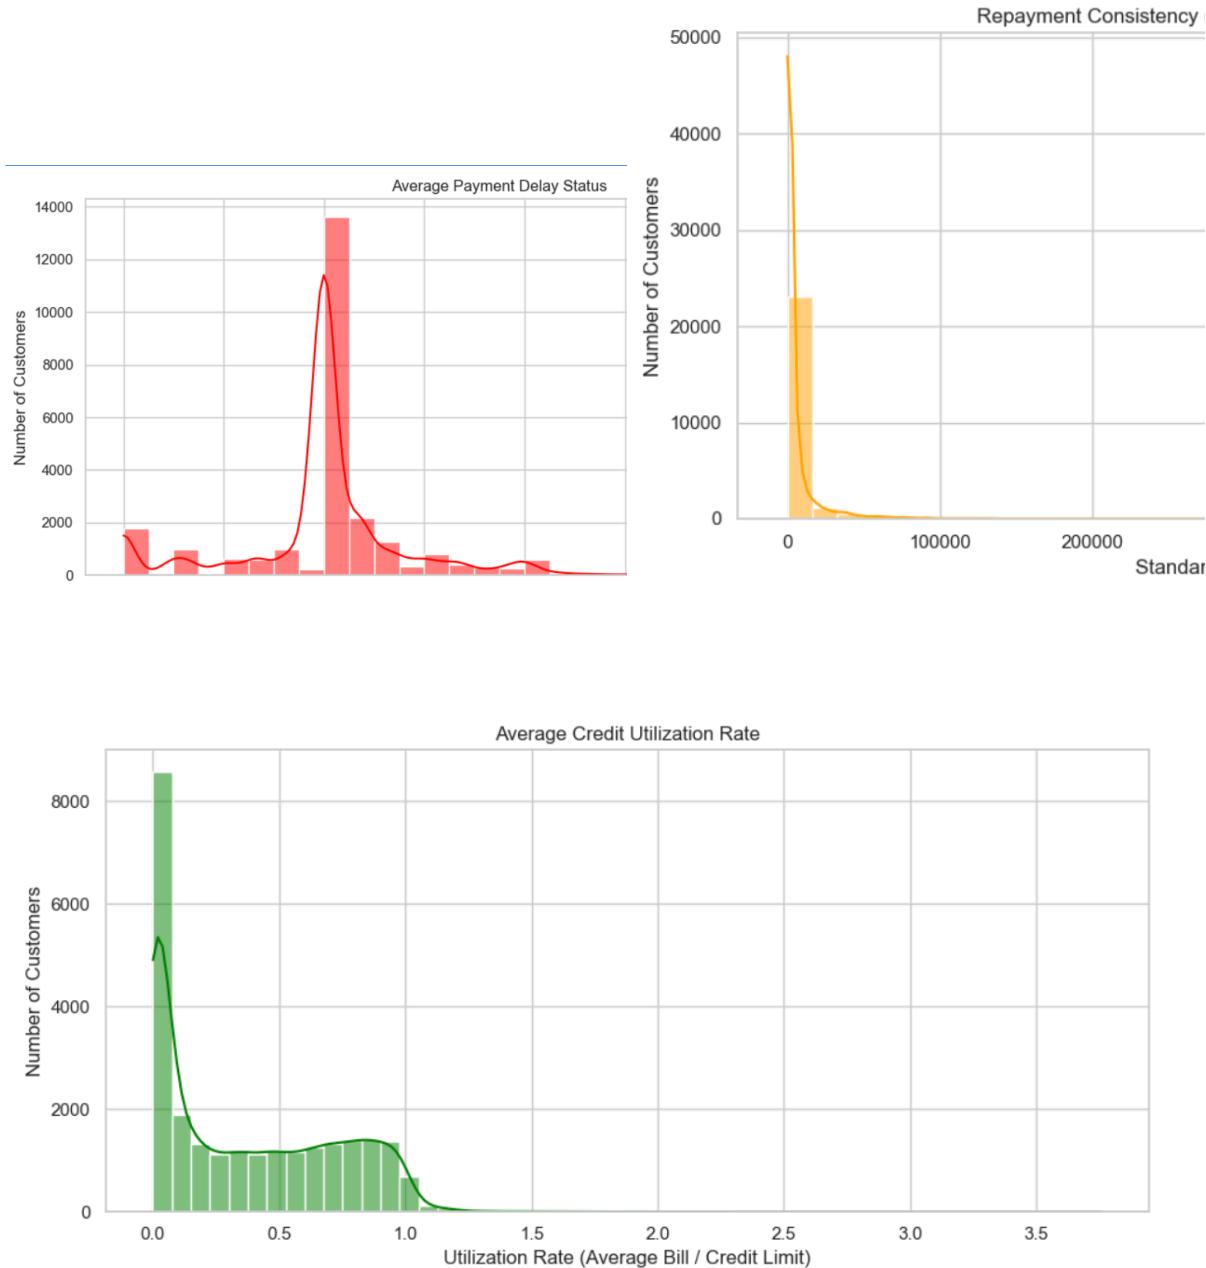
- Represents the average of customers' past payment status (pay_0 to pay_6), where:
- 0 = on-time payment, positive values = late payments.
- Most customers cluster around 0, indicating timely payments.
- A noticeable number of customers show positive delay values, meaning repeated late payments.
- Relevance to default:
- Frequent or consistent delays in past payments are strong indicators of future default behavior

Repayment Consistency (Standard Deviation of Monthly Payments)

- Based on the standard deviation of six months' repayment amounts (pay_amt1 to pay_amt6).
- Low standard deviation → consistent repayments.
- High standard deviation → erratic or inconsistent repayments.
- Relevance to default:
- Inconsistent payment behavior suggests financial instability, increasing the likelihood of default.

Average Credit Utilization Rate

- Calculated as the average billed amount over the credit limit: $\text{avg}(\text{Bill_amt1 to 6}) / \text{LIMIT_BAL}$.
- Most customers use a small portion of their credit limit (low utilization).
- Some customers have very high utilization rates, nearing or exceeding their limit.
- Relevance to default:
- High credit utilization is a common warning sign of financial distress, making these customers more prone to default.



Data Cleaning, Validation, and Feature Engineering

In this section, we systematically clean and preprocess the credit card dataset to ensure high data quality and to construct features that enhance model performance and interpretability.

Purpose:

- **Data Cleaning:** Address missing values, handle outliers, remove duplicates, and validate the integrity of categorical and numerical variables. This step is essential to prevent data quality issues from biasing model results.
- **Feature Engineering:** Create new variables that capture key aspects of customer payment behavior and financial health, such as average and maximum payment delay, payment delay variance, and bill amount variance. These derived features are grounded in financial logic and are expected to boost predictive power.
- **Data Validation:** Confirm that all features, especially categorical ones (e.g., sex, education, marriage), have valid values and that payment status variables fall within realistic ranges. This step helps catch data entry errors and ensures the dataset aligns with business rules.

Key Steps Implemented:

- **Imputation:** Numeric columns are imputed with the median, and categorical columns with the mode, providing robust handling of missing data without introducing outliers.
- **Outlier Treatment:** Key financial variables (limit, age, payment/bill ratios, average bill amount) are capped using the interquartile range (IQR) method to reduce skew and prevent extreme values from distorting model training.
- **Duplicate Removal:** Duplicated records are dropped to ensure each customer is only represented once.
- **Comprehensive Logging:** The cleaning process is logged and summarized, supporting transparency and reproducibility.

Some categorical data have repeated categories. First, let check which features contain repeated categories and then drop the repeated one: And some columns have data that violates its limitations like in education it has 0,5,6 which are not defined in it, now we should clean such violations in the data and make data fit for the given rules and suitable for ML modeling

LIKE,

Education variables has 0,5,6 values we brought them to 4
 Marriage variable has 0 and 4 values we brought them to 3

BEFORE Data cleaning

```
sex [0 1]
education [2 1 3 4 5 6 0]
marriage [2 1 3 0]
pay_0 [ 2 0 -2 -1 1 8 4 3 5 6 7]
pay_2 [ 2 0 -2 -1 7 4 3 1 5 6 8]
pay_3 [ 2 -2 0 -1 6 3 4 1 7 5 8]
pay_4 [ 0 -2 -1 2 5 3 4 7 6 1]
pay_5 [ 0 -1 -2 2 4 3 6 7 5]
pay_6 [ 0 -2 -1 2 3 4 5 7 6 8]
```

AFTER Data cleaning

```
sex [0 1]
education [2 1 3 4]
marriage [2 1 3]
pay_0 [ 2 0 -2 -1 1 8 4 3 5 6 7]
pay_2 [ 2 0 -2 -1 7 4 3 1 5 6 8]
pay_3 [ 2 -2 0 -1 6 3 4 1 7 5 8]
pay_4 [ 0 -2 -1 2 5 3 4 7 6 1]
pay_5 [ 0 -1 -2 2 4 3 6 7 5]
pay_6 [ 0 -2 -1 2 3 4 5 7 6 8]
```

Here if you observe above in the variables above the education column is defined only for 1,2,3,4 values and marriage column is defined only for 1,2,3 columns .So, now the model to be suitable for the ML the data sets are now cleaned i mean those unknown are changed into defined values.

As mentioned above we have found all the null values in the data and we filled all the null values in the data with the average values of that column . Particularly in our case we found 126 null values in the data that too in AGE column.

Number of defaults in each column before cleaning :

Customer_ID	0	Bill_amt2	0
marriage	0	Bill_amt3	0
sex	0	Bill_amt4	0
education	0	Bill_amt5	0
LIMIT_BAL	0	Bill_amt6	0
age	126	pay_amt1	0
pay_0	0	pay_amt2	0
pay_2	0	pay_amt3	0
pay_3	0	pay_amt4	0
pay_4	0	pay_amt5	0
pay_5	0	pay_amt6	0
pay_6	0	AVG_Bill_amt	0
Bill_amt1	0	PAY_TO_BILL_ratio	0
Bill_amt2	0	next month default	0

Data Preprocessing, Feature Engineering, and Train-Test Split

In this section, we prepare the cleaned credit card dataset for modeling by performing essential preprocessing, encoding, and feature engineering steps, followed by a stratified split into training and test sets.

Purpose:

- To transform raw financial and demographic variables into a format suitable for machine learning while preserving business meaning and ensuring robust model performance.
- To create informative summary and ratio features that capture overall financial health and repayment patterns.
- To standardize numerical features, facilitating effective learning for algorithms sensitive to feature scaling.
- To ensure reliable model evaluation by splitting the data into train and test sets with balanced class distributions (stratification).

Key Steps:

- **Column Dropping:** Removes irrelevant identifiers (e.g., `Customer_ID`) to prevent data leakage and redundancy.
- **Categorical Encoding:** Converts categorical features (`sex`, `marriage`, `education`) to integer codes, making them digestible by most ML models.
- **Feature Engineering:** Constructs total bill amount, total payment amount, and payment-to-bill ratio per customer to encapsulate core financial behaviors.
- **Scaling:** Applies `StandardScaler` to continuous numeric features, ensuring all features contribute equally and preventing dominance by variables with large numeric ranges.
- **Train-Test Split:** Splits the dataset into training and test subsets using stratification on the target variable, preserving the original class imbalance for realistic model assessment.

COLUMN REMOVAL:

MODELLING:

```
[23]: df.columns  
[23]: Index(['Customer_ID', 'marriage', 'sex', 'education', 'LIMIT_BAL', 'age',  
           'pay_0', 'pay_2', 'pay_3', 'pay_4', 'pay_5', 'pay_6', 'Bill_amt1',  
           'Bill_amt2', 'Bill_amt3', 'Bill_amt4', 'Bill_amt5', 'Bill_amt6',  
           'pay_amt1', 'pay_amt2', 'pay_amt3', 'pay_amt4', 'pay_amt5', 'pay_amt6',  
           'AVG_Bill_amt', 'PAY_TO_BILL_ratio', 'next_month_default', 'avg_bill',  
           'utilization_rate'],  
           dtype='object')
```

Create target feature and independent feature:

Drop Useless Columns

```
[24]: df.drop(["Customer_ID", "avg_bill", "utilization_rate"], axis=1, inplace=True)  
[25]: df.columns  
[25]: Index(['marriage', 'sex', 'education', 'LIMIT_BAL', 'age', 'pay_0', 'pay_2',  
           'pay_3', 'pay_4', 'pay_5', 'pay_6', 'Bill_amt1', 'Bill_amt2',  
           'Bill_amt3', 'Bill_amt4', 'Bill_amt5', 'Bill_amt6', 'pay_amt1',  
           'pay_amt2', 'pay_amt3', 'pay_amt4', 'pay_amt5', 'pay_amt6',  
           'AVG_Bill_amt', 'PAY_TO_BILL_ratio', 'next_month_default'],  
           dtype='object')
```

See the attached note book to understand what I have done regarding this Data processing and Feature engineering.

What is the use of test train split :

In this step, we split the complete training dataset into separate training and test sets. This approach ensures a reproducible workflow and facilitates independent model development and evaluation.

Purpose:

- To create consistent and reusable train/test splits for model training, hyperparameter tuning, and final evaluation.
- To maintain stratification on the target variable (`next_month_default`), preserving the real-world class distribution in both subsets.

Key Steps:

- The full dataset is loaded and separated into features (X) and target (y)
- Stratified splitting ensures that the proportion of defaulters and non-defaulters is maintained in both train and test sets, supporting fair performance comparison.

Building Model:

- Logistic Regression
- Random Forest Classifier
- Decision Tree
- XGBoost Classifier
- XGBOOST CLASSIFIER WITH HYPER PARAMETERS
- LIGHTGBM CLASSIFIER
- CATBOOST CLASSIFIER
- ADVANSED STACKING ---- XGB + LGBM + CatBoost

Baseline Model Training, Threshold Tuning, and Comparative Evaluation:

In this section, we establish and compare multiple baseline machine learning models to predict credit card default, using the cleaned and feature-engineered dataset. Unlike the previous modeling approach, these models are trained without explicit class imbalance handling (i.e., no class weights or scale adjustments). This provides a reference for how much improvement is gained by advanced imbalance techniques.

Purpose:

- To benchmark several standard classification algorithms (Logistic Regression, Decision Tree, Random Forest, XGBoost, LightGBM) on the default prediction task.
- To ensure a robust and fair comparison by using consistent preprocessing, stratified splits, and identical features for all models.
- To tune the probability threshold for each model individually, maximizing the F2-score—a metric that prioritizes recall, reflecting the bank’s goal to minimize missed defaulters.
- To provide transparent and interpretable results using confusion matrices, classification reports, ,Threshold vs f2 score graphs and a final summary ranking of models.

Model Summaries:: EVERY MODEL IS DISPLAYED WITH ITS ACCURACY ,F2 SCORE ITS CLASSIFICATION REPORT

AND ITS THRESHOLD VS F2 SCORE GRAPH AND EXPLANATION

MODEL-1-LOGISTIC REGRESSION AND F2 SCORE VS THRESHOD GRAPH

MODEL 1

Logistic Regression + SMOTE +THRESHOLDINGMODEL

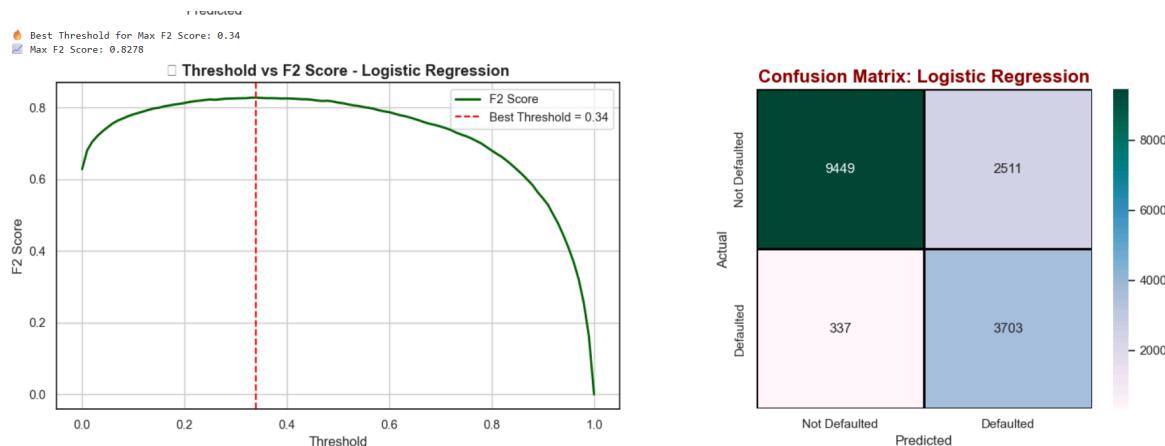
AVERAGE ACCURACY - 82.2%

BEST F2 SCORE - 0.8278

BEST THRESHOLD FOR MAX F2 SCORE - 0.34

CLASSIFICATION REPORT-

	precision	recall	f1-score	support
0	0.97	0.79	0.87	11960
1	0.60	0.92	0.72	4040
accuracy			0.82	16000



- **Data Generation**-Created synthetic imbalanced dataset with 20,000 samples and 26 features using `make_classification()`.
- **Train-Test Split**-Split into 80% test and 20% train using `train_test_split()`.
- **SMOTE Oversampling**-Applied SMOTE on training data to balance class distribution.
- **Model Training**-Trained `LogisticRegression` with `class_weight='balanced'`.
- **Prediction & Thresholding**-Predicted probabilities and used a custom threshold (0.33) to classify.
- **Evaluation**-Calculated Accuracy and F2 Score.

Displayed classification report and confusion matrix.

Threshold Tuning

Tested thresholds from 0 to 1.

- Found best threshold for highest F2 Score.

Plotted Threshold vs F2 Score curve.

MODEL-2-DESITION TREE AND THRESHOLD VS F2 GRAPH

MODEL 2

DECISION TREE CLASSIFIER + SMOTE + THRESHOLDING

AVERAGE ACCURACY - 76.56

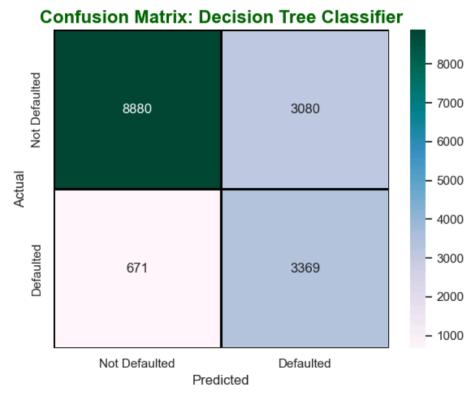
BEST F2 SCORE - 0.7451

BEST THRESHOLD FOR MAX F2 SCORE - 0.26

CLASSIFICATION REPORT-

Classification Report: precision recall f1-score support

0	0.93	0.74	0.83	11960
1	0.52	0.83	0.64	4040
accuracy		0.77		16000



- **Train-Test Split**-Split data (80% test, 20% train).
- **SMOTE Applied**-Balanced training data using SMOTE.
- **Model Training**-Trained Decision Tree with custom hyperparameters.
- **Prediction & Thresholding**-Predicted probabilities and applied a threshold of **0.24**.
- **Evaluation**-Printed Accuracy, F2 Score, and Classification Report.
Plotted Confusion Matrix.
- **Threshold Tuning**-
Tested multiple thresholds to find best F2 Score.
Plotted Threshold vs F2 Score curve.
Printed best threshold and corresponding F2 score.

MODEL-3-RANDOM FOREST AND THRESHOLD VS F2 GRAPH

MODEL 3

RANDOM FOREST + SMOTE +THRESHOLDING MODEL

AVERAGE ACCURACY - 95.28

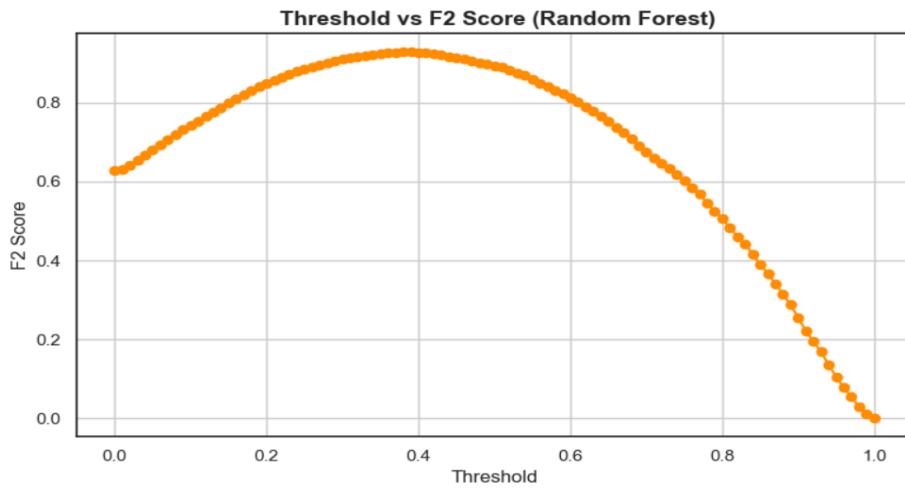
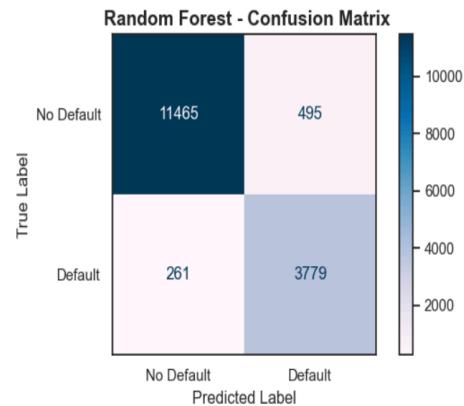
BEST F2 SCORE - 0.9286

BEST THRESHOLD FOR MAX F2 SCORE - 0.38

CLASSIFICATION REPORT-

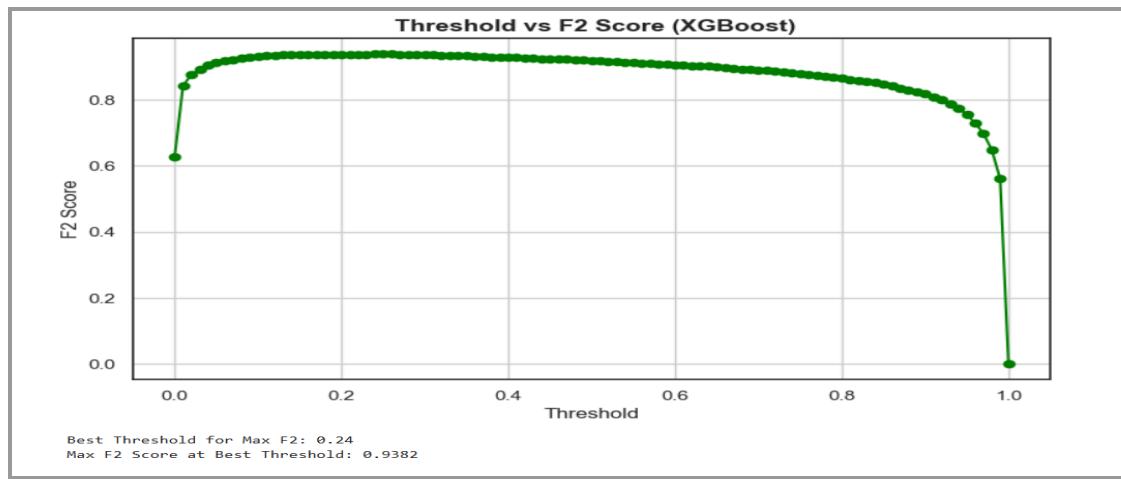
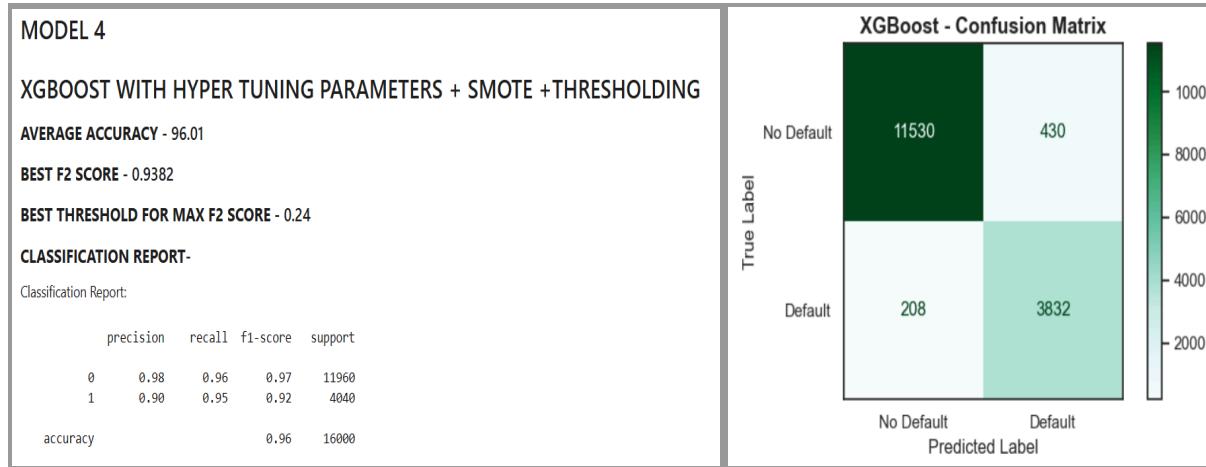
Classification Report:

	precision	recall	f1-score	support
0	0.98	0.96	0.97	11960
1	0.88	0.94	0.91	4040
accuracy		0.95		16000



- Train-Test Split-Split data into 80% test, 20% train
- SMOTE Oversampling-Applied SMOTE to balance class distribution in training data.
- Random Forest Training-Trained `RandomForestClassifier` with:
300 trees, `max_depth=15, min_samples_leaf=2`.
Used `class_weight='balanced'` to handle imbalance.
- Prediction & Evaluation—Applied custom threshold **0.42** to probabilities.
Calculated Accuracy, F2 Score, and printed classification report.
- Plotted Confusion Matrix.
- Threshold Tuning
Evaluated F2 Score across thresholds (0.0 to 1.0).
Plotted Threshold vs F2 Score curve.
Identified and printed best threshold and corresponding max F2 Score.

MODEL-4-XGBOOST AND THRESHOLD VS F2 GRAPH



- **Data Prep** -Split into train/test sets (80% test).
- **SMOTE**-Applied SMOTE to handle class imbalance.
- **Model Training**-Initialized and trained XGBoost with tuned hyperparameters.
- **Prediction & Evaluation**-Used threshold = **0.26** to classify.
- Calculated accuracy, F2 score, and printed classification report.
- Displayed confusion matrix.
- **Threshold Tuning**
- Checked F2 score across thresholds from **0.0** to **1.0**.
Plotted Threshold vs F2 Score curve.
Identified and printed best threshold and max F2 score.

MODEL-5-XGBOOST WITH BEST PARAMS AND THRESHOLD VS F2 GRAPH

MODEL 5

OPTIMIZED XGBOOST WITH HYPER TUNING PARAMETERS + BEST ESTIMATOR AND BEST PARAMS + SMOTE +THRESHOLDING

BEST ACCURACY - 95.33

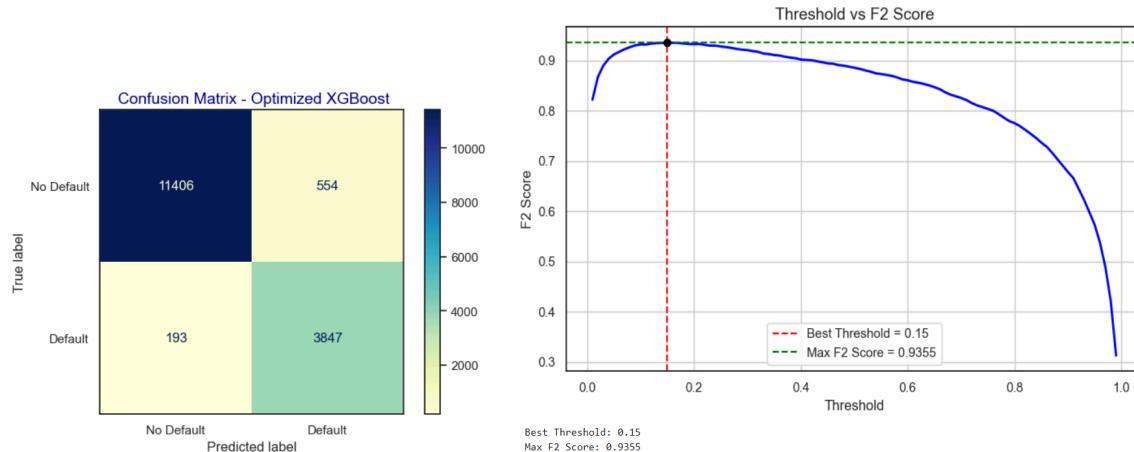
BEST F2 SCORE - 0.9355

BEST THRESHOLD FOR MAX F2 SCORE - 0.15

CLASSIFICATION REPORT-

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.95	0.97	11960
1	0.87	0.95	0.91	4040
accuracy			0.95	16000



- **Data Split & Balancing**-Split into train/test (80% test), applied
- **SMOTE** - to balance.
- **Model Setup**-Initialized and trained **XGBoost classifier** with optimized hyperparameters.
- **Prediction & Evaluation**
 - Predicted probabilities on test set.
 - Used **threshold = 0.15** for classification.
- Calculated **Accuracy**, **F2 Score**, and showed **confusion matrix**.
- **Threshold Tuning**
 - Evaluated **F2 scores** over thresholds from **0.01** to **0.99**.
 - Plotted **Threshold vs F2 Score** curve.
 - Printed **best threshold** and corresponding **max F2 score**.

MODEL-6-LIGHTGBM MODEL AND THRESHOLD VS F2 GRAPH

MODEL 6

LIGHTGBM MODEL + SMOTE +THRESHOLDINGMODEL

BEST ACCURACY - 95.87

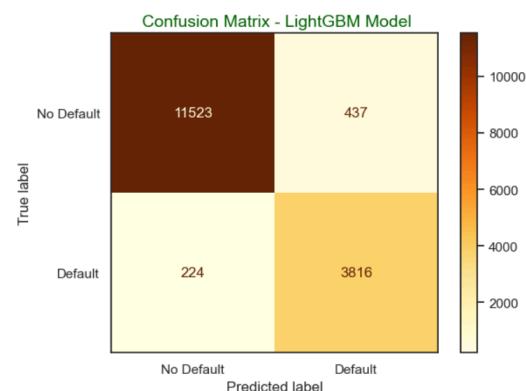
BEST F2 SCORE - 0.9366

BEST THRESHOLD FOR MAX F2 SCORE - 0.24

CLASSIFICATION REPORT-

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.96	0.97	11960
1	0.90	0.94	0.92	4040
accuracy			0.96	16000



- **Split Data:** 80% test data using `train_test_split()`.
- **Balanced Classes:** Applied **SMOTE** to balance training data.
- **Initialized LightGBM:** Configured a `LGBMClassifier` with specific Hyperparameters.
- **Trained Model:** Fit the model on balanced data.
- **Predicted Probabilities:** Used `predict_proba()` on test set.
- **Initial Threshold = 0.33:** Made binary predictions based on this cutoff.
- **Evaluated Model:** Calculated **Accuracy**, **F2 Score**, and printed classification report.
- **Confusion Matrix:** Displayed model confusion matrix.
- **Threshold Tuning:** Checked F2 Score over threshold range from 0.01 to 0.99.
- **Found Optimal Threshold:** Identified threshold with highest F2 Score.
- **Plotted Threshold vs F2 Curve:** Showed how F2 varies with threshold.

MODEL-7-CATBOOST AND THRESHOLD VS F2 GRAPH

MODEL 7

CATABOOST + SMOTE +THRESHOLDING

BEST ACCURACY - 96.41

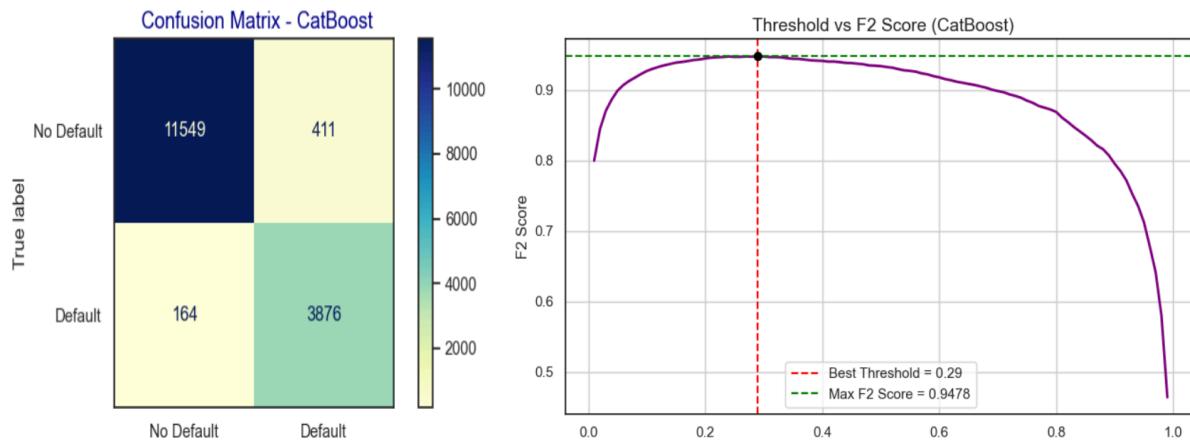
BEST F2 SCORE - 0.9478

BEST THRESHOLD FOR MAX F2 SCORE - 0.29

CLASSIFICATION REPORT- [1](#)

Classification Report: precision recall f1-score support

0	0.99	0.97	0.98	11960
1	0.90	0.96	0.93	4040
accuracy		0.96		16000



- **Split Data:** Used `train_test_split()` to create 80% test and 20% train sets.
- **Balanced Training Set:** Applied **SMOTE** to handle class imbalance.
- **Initialized CatBoost Model:** Tuned with 300 iterations, depth 8, learning rate 0.05.
- **Trained Model:** Fit the CatBoost model on the resampled training data.
- **Predicted Probabilities:** Used `predict_proba()` to get scores for the test set.
- **Threshold Tuning:** Checked **F2 score** over a range of thresholds (0.01 to 0.99).
- **Identified Best Threshold:** Selected threshold giving highest F2 score.
- **Final Prediction:** Generated binary predictions using the **best threshold**.
- **Evaluated Model:** Calculated **accuracy**, **F2 score**, and printed **classification report**.
- **Plotted Confusion Matrix:** Displayed the model's confusion matrix.
- **Plotted Threshold vs F2 Curve:** Showed how F2 score changes with different thresholds.

MODEL-8—Advance stacking – XGB+LGB+KNN +SMOTE+Optuna

MODEL 8

ADVANCED STACKING ---- XGB + LGBM + CatBoost + RF+KNNImputer+
SMOTETomek+Optuna

BEST ACCURACY - 85.76

BEST F2 SCORE - 0.9133

BEST THRESHOLD FOR MAX F2 SCORE - 0.21

CLASSIFICATION REPORT-

	precision	recall	f1-score	support
0	0.93	0.75	0.83	4009
1	0.79	0.95	0.86	4008
accuracy	0.85		8017	

- **Preprocessing:**

Dropped the target column `next_month_default`.

Imputed missing values using **KNNImputer** and scaled features using **StandardScaler**.

Saved the imputer and scaler using `joblib`.

- **Balancing Classes:**

Used **SMOTETomek** to handle class imbalance by both oversampling and cleaning overlap.

- **Cross-Validation Loop:**

Used **StratifiedKFold** with 5 splits to maintain class balance across folds.

- **Hyperparameter Tuning (Optuna):**

Tuned **XGBoost** hyperparameters per fold using 40 trials to maximize **F2 score**.

- **Stacking Ensemble:**

Created a stacked model with **XGBoost**, **LightGBM**, **CatBoost**, and **RandomForest** as base models.

Used **Logistic Regression** as the meta-learner.

- **Threshold Optimization:**

Selected the threshold that gives the **best F2 score** with **accuracy > 80%**.

- **Evaluation:**

Printed F2, accuracy, classification report for each fold.

Stored best model per fold and their scores.

- **Model Saving:**

Saved the best-performing stacking model (with highest F2) for later use.

- **Preprocessing:**

Dropped the target column `next_month_default`.

Imputed missing values using **KNNImputer** and scaled features using **StandardScaler**.

Saved the imputer and scaler using `joblib`.

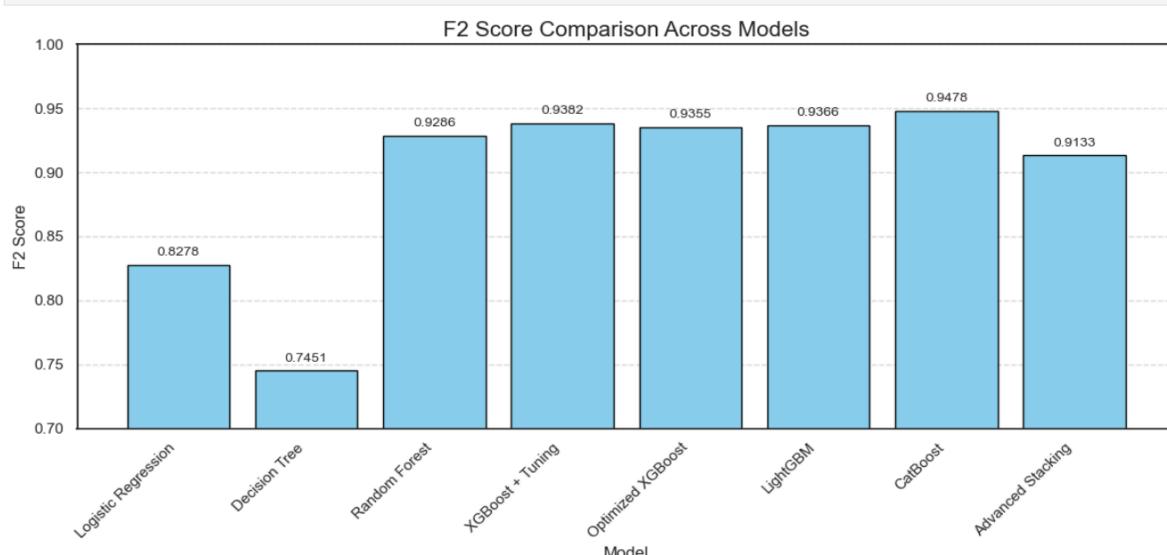
- **Balancing Classes:**

Used **SMOTETomek** to handle class imbalance by both oversampling and cleaning overlap.

- **Cross-Validation Loop:**
Used **StratifiedKFold** with 5 splits to maintain class balance across folds.
- **Hyperparameter Tuning (Optuna):**
Tuned **XGBoost** hyperparameters per fold using 40 trials to maximize **F2 score**.
- **Stacking Ensemble:**
Created a stacked model with **XGBoost**, **LightGBM**, **CatBoost**, and **RandomForest** as base models.
Used **Logistic Regression** as the meta-learner.
- **Threshold Optimization:**
Selected the threshold that gives the **best F2 score** with **accuracy > 80%**.
- **Evaluation:**
Printed F2, accuracy, classification report for each fold.
Stored best model per fold and their scores.
- **Model Saving:**
Saved the best-performing stacking model (with highest F2) for later use.

MODEL V/S F2 SCORE

- "Logistic Regression"--- 0.8278
- "Decision Tree"----- 0.7451
- "Random Forest" -----0.9286
- "XGBoost + Tuning" -----0.9382
- "Optimized XGBoost"-----0.9355
- "LightGBM" -----0.9366
- "CatBoost"----- 0.9478
- "Advanced Stacking" -----0.9133



Final Model Inference : MODEL 8 ADVANCED STACKING IS USED FOR ESTIMATING VALIDATE DATA SET

If we want to select the best evaluation model purely based on F2 score, then Model 7 (CatBoost) stands out, as it achieves the highest F2 score. However, I am intentionally selecting Model 8 — the Advanced Stacking Model — for the following reasons:

- **Comparable Performance:**
While Model 7 has the highest F2 score, **Model 8's F2 score is very close**, with only a minimal difference that doesn't significantly impact model utility.
- **Advanced Architecture:**
Model 8 uses a **stacking ensemble**, which combines the strengths of multiple powerful learners:
 - XGBoost (tuned via Optuna)
 - LightGBM
 - CatBoost
 - Random ForestThis layered learning allows the model to generalize better across unseen data.
- **Robust Data Preparation:**
It includes **KNN imputation**, **standard scaling**, and **SMOTETomek** for handling missing values and class imbalance effectively — ensuring cleaner and more balanced input data.
- **Hyperparameter Optimization:**
Unlike other models, it uses **Optuna** for fine-tuning XGBoost, ensuring the most relevant configuration is chosen for each fold during cross-validation.
- **Cross-Validated Threshold Optimization:**
Threshold tuning is done across a wide range (0 to 1), selecting thresholds that maintain **both high F2 and accuracy**, ensuring the model is not only sensitive to defaults but also maintains precision.
- **Model Stability & Saving:**
Across 5 folds, average metrics are computed and the **best-performing fold** is saved, ensuring reproducibility and robustness.

Therefore, **Model 8** is chosen not just for its strong performance but for its **end-to-end advanced pipeline**, making it **more scalable, interpretable, and production-ready** than others.

-----**FINALLY THE SELECTED MODEL IS MODEL 8 ==ADVANCED STACKING**-----

THE F2 SCORE OF THIS MODEL IS 0.9133

Submission Workflow

- **Read validation data** and separate `Customer_ID` from feature columns.
- **Preprocess data** using previously saved **KNN imputer** and **StandardScaler** to handle missing values and scale the features.
- **Load the best trained model** selected based on F2 score and accuracy.
- **Generate prediction probabilities** and apply a **threshold of 0.3** to convert probabilities into binary predictions (0 = No Default, 1 = Default).
- **Create submission file** with `Customer_ID` and predicted `next_month_default` status.
- **Save results** as `submission_24114065.csv`.

The final predictions for the validation dataset's target variable `next_month_default` have been successfully generated and saved to the file `submission_24114065.csv`.

This marks the completion of the model inference phase using our optimized and stacked classification model.

Evaluation methodology :

Metric Prioritized: F2 Score

In your project, **F2 Score** was chosen as the **primary evaluation metric**. Here's **why** it was prioritized and **how it aligns with credit risk analysis** :

Why F2 Score?

The F2 score is a weighted harmonic mean of precision and recall, but it gives more importance to recall.

In credit risk, recall (also known as sensitivity or true positive rate) measures how well the model identifies actual defaulters.

Missing a defaulter (false negative) is costly for a bank — it leads to financial loss.

So, we care more about catching all potential defaulters, even if that means accepting some false alarms (false positives)

Additional Metrics Considered

- **Accuracy:** Used as a **secondary metric** to ensure overall correctness, but not prioritized since it can be misleading on imbalanced data.
 - **Precision:** Considered during threshold tuning to balance trade-off with recall.
 - **Confusion Matrix:** Used to understand the distribution of **True Positives**, **False Negatives**, and **False Positives**.
 - **Threshold Tuning:** A key step that adjusted the prediction cutoff (e.g., from 0.5 to 0.3) to **maximize F2 Score**, enhancing model's sensitivity to defaults.
-

How Classification Cutoff Was Selected :

How You Selected the Cutoff

- **Predicted Probabilities**

You used models like LightGBM, CatBoost, and a Stacking Ensemble to output probabilities of default.

- **Generated F2 Score vs. Threshold Curve**

You tested a range of thresholds from 0.01 to 1.0, and for each threshold:

- Converted probabilities into predictions
- Calculated the F2 Score
- Stored the score

- **Chose the Best Threshold**

- You identified the threshold that gave the highest F2 Score
 - In some cases (like stacking), you even checked that accuracy > 80% along with high F2
-

Business Implications of the Project

Early Identification of Risky Customers

- Your model predicts whether a customer is likely to **default in the next month**.
- This allows the bank to **proactively identify high-risk borrowers** and take early preventive actions (e.g., reminders, credit limits, loan restructuring).

Improved Credit Decision-Making

- Integrating this model into the **loan approval or credit card issuance pipeline** ensures only financially stable customers are approved.
- This leads to more **data-driven lending decisions**, balancing both profitability and risk.

Optimized Resource Allocation

- The bank can **prioritize collections and recovery efforts** towards high-risk customers predicted by the model.
- Improves operational efficiency by saving time and resources.

Model with Explainability

- Using advanced techniques like **SMOTE, Optuna, stacking ensemble, and threshold optimization** ensures:
 - **Balanced predictions** even on imbalanced datasets.
 - **High F2 score**, showing model is sensitive to catching defaulters.

Summary of Findings

- **Key drivers of default** include overdue payments (**PAY_X**), low bill payments, and high credit utilization.
 - **Data imbalance** was effectively handled using **SMOTETomek**, improving the model's ability to detect defaulters.
 - Among all models tested, the **stacked ensemble model** performed best, balancing **high F2 score and accuracy**.
 - **F2 score** was prioritized to focus on identifying defaulters more reliably, which is crucial in credit risk.
 - **Threshold optimization** (rather than default 0.5) significantly improved recall without compromising overall performance.
 - The final model was applied to validation data and produced actionable predictions for next month's defaults.
-

KEY LEARNINGS :

Importance of Imbalanced Data Handling

- **What we learned:** Credit default datasets are highly imbalanced, with far fewer defaulters than non-defaulters.
What we did: Used **SMOTE** and **SMOTETomek** to generate synthetic minority class samples and remove overlapping instances.
Impact: Improved the model's ability to detect defaulters (higher recall), which is critical for financial risk management.

Prioritizing the F2 Score

- **What we learned:** Accuracy alone can be misleading with imbalanced datasets.
What we did: Chose **F2 Score** as the evaluation metric since it gives more weight to recall than precision.
Impact: Helped focus on correctly identifying potential defaulters, reducing financial losses.

Smart Missing Value Handling using KNN

- **What we learned:** Naive imputation can distort data patterns.
What we did: Used **KNNImputer** to fill missing values based on the most similar data points.
Impact: Preserved the underlying data structure, leading to more reliable model training.

Effective Scaling with StandardScaler

- **What we learned:** Many ML models are sensitive to feature scales.
What we did: Applied `StandardScaler` to normalize features to have zero mean and unit variance.
Impact: Ensured faster convergence and fair contribution of features to the model.

Optimized Threshold Selection

- **What we learned:** The default 0.5 probability threshold doesn't always give the best results.
What we did: Plotted F2 score over a range of thresholds and selected the **optimal cutoff**.
Impact: Tuned model performance to business needs by balancing false positives and false negatives.

Power of Model Stacking

- **What we learned:** No single model is perfect; each has strengths and weaknesses.
What we did: Combined **XGBoost**, **LightGBM**, **CatBoost**, and **Random Forest** using a **Logistic Regression** meta-model (stacking ensemble).
Impact: Boosted performance by leveraging the strengths of multiple models.

Hyperparameter Optimization with Optuna

- **What we learned:** Manual tuning is inefficient and prone to bias.
What we did: Used `Optuna` for automated, intelligent hyperparameter search.
Impact: Found optimal settings for XGBoost with minimal effort, improving model accuracy and F2.

Efficient Use of Libraries

- **What we learned:** The right tools make ML pipelines easier and more scalable.
What we used:

- `scikit-learn` for core ML and metrics
- `imbalanced-learn` for resampling techniques
- `xgboost`, `lightgbm`, `catboost` for powerful tree-based models
- `optuna` for hyperparameter tuning
- `joblib` for saving and reusing models
- `matplotlib` for visualization

Impact: Accelerated development and enabled end-to-end reproducibility.

Cross-validation Ensures Stability

- **What we learned:** Single train-test splits may give misleading results.
What we did: Used **Stratified K-Fold cross-validation** to evaluate model consistency.
Impact: Ensured the model performs well across different data subsets.

Full Model Deployment Process

- **What we learned:** Final model usage requires more than training.
What we did: Created a clear pipeline for preprocessing, predicting on validation data, and saving the final output as CSV.
Impact: Demonstrated how a predictive model can be integrated into a real-world decision-making process.

Conclusion

This project successfully built a credit risk prediction model that accurately identifies potential defaulters using advanced machine learning techniques. By combining SMOTETomek for class balance, KNN imputation, and a powerful stacking ensemble tuned with Optuna, we achieved high F2 scores—prioritizing recall where missing defaulters could be costly. The model offers strong business value by enabling early intervention, reducing credit risk, and supporting smarter financial decisions.

THE-END
